Membership is Ownership: A Robust Model Watermark Detection with Quantile Regression

Anonymous Author(s)
Affiliation
Address
email

Abstract

Large-scale diffusion models are valuable intellectual property, yet can be stolen and adapted via downstream fine-tuning. Existing diffusion watermarking embeds explicit signatures into parameters or behavior, which may degrade utility and can be removed by post-hoc fine-tuning. We propose Membership Is Ownership, an ownership verification framework that treats a private member evidence set as a watermark and leverages the model's intrinsic memorization. Given a suspect model, we compute a multi-timestep *t-error* score by forward noising and single-step reconstruction, then aggregate with a lower quantile (Q25). To control false-positive rates, we fit a conditional Gaussian model in log-score space and obtain closed-form quantile thresholds at arbitrary target FPR. The protocol compares the suspect model against public baselines via hypothesis tests and effect sizes, producing statistical evidence rather than a single score. On CIFAR-10/100, STL-10, and CelebA, all datasets achieve $|d| > 18$ against public baselines; these separations persist after MMD fine-tuning (simulating post-theft adaptation), provided the adversary does not retrain from scratch.

1 Introduction

Large-scale generative diffusion models have fueled numerous profitable downstream applications. To obtain such powerful models, a huge amount of computation and data resources are consumed, which renders the trained diffusion models valuable intellectual property (IP) for tech companies and individuals. However, these IP assets are vulnerable to various unauthorized uses by adversaries seeking to leverage these models for customized, usually commercial applications.

In practice, an adversary can obtain a checkpoint via model extraction attacks or illicit transfer, then fine-tune the model to a profitable domain or product requirement Wang et al. [2025], Kim et al. [2024]. Such post-theft fine-tuning undermines the model owner's copyright and poses significant challenges for reliable ownership verification. To make a strong claim, the owner must produce evidence that is not only accurate under stringent false-positive budgets, and also valid after realistic fine-tune or adaptation.

Most existing diffusion IP protection methods embed explicit signatures into model parameters or output behavior [Peng et al., 2023, Zhao et al., 2023, Fernandez et al., 2023, Liu et al., 2023]. While effective in controlled settings, such watermarks may impose utility costs and can be weakened by fine-tuning because the model thief can use every means to remove such explicit watermarks. Recently, an alternative watermark method has emerged, which is to leverage the model's intrinsic memorization [Carlini et al., 2023, Somepalli et al., 2023, Gu et al., 2025] , where the reconstruction error of diffusion models on training examples is systematically lower than on non-training samples. However, as the models' generalization capability increases, the difference between training samples and non-training samples becomes marginal, resulting in the "low error-based" method

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

being insufficient for an ownership claim. Moreover, most modern diffusion models may also use substantially overlapped training data from the public Internet, which increases the difficulty of accurate attribution.

These practical challenges necessitate a robust watermark detection method that can distinguish the IP-protected model from public reference models under a statistically controlled decision rule. To fill this research gap, in this work, we propose **Membership is Ownership**, an ownership verification framework that can accurately detect and verify whether a suspect model in the thief's possession is illegally obtained and fine-tuned from the model owner's base model, using a designated owner signature. At the heart of our approach lies a simple but principled observation: a model retains systematic memory of its training data, manifesting as lower reconstruction error on seen examples. We exploit this signal by constructing a private evidence set held by the model owner, which serves as a persistent watermark without modifying the model itself. To extract a robust signal from a suspect model, we introduce a multi-timestep *t-error* score—computed via forward noising followed by single-step reconstruction—and summarize it through a lower quantile (`Q25`) to suppress noise while preserving the memorization signature. Rather than reporting a single black-box decision, we anchor the verification protocol on a principled statistical foundation: a conditional Gaussian model fitted in log-score space yields closed-form quantile thresholds at arbitrary false-positive rates, enabling transparent and auditable comparisons against public baselines through hypothesis tests and effect sizes. The resulting separations are striking. Against public baselines, we observe $|d| > 18$, margins that remain robust even under MMD fine-tuning, which simulates realistic post-theft adaptation, provided the adversary stops short of full retraining or targeted memorization erasure. The paper unfolds the following contributions.

- **Conceptual.** We reframe ownership verification from single-sample membership inference into a population-level hypothesis-testing problem, using a private member evidence set held by the owner as a non-invasive watermark that leaves model parameters and output quality intact.
- **Methodological.** We combine multi-timestep `Q25` aggregation with a Gaussian quantile regression in log-score space to yield closed-form FPR control, enabling auditable verification at arbitrary significance levels without retraining.
- **Empirical.** We validate the framework on DDIM (4 datasets) and Stable Diffusion v1.4 (LoRA r=$64/256$, full fine-tuning); the population-level separation persists under MMD, SGD, weight-noise, and PEFT fine-tuning, with all DDIM-from-scratch settings achieving $|d| > 18$ and the SD-PEFT regime passing the paired-delta protocol. MiO matches WDM and Zhao's post-theft robustness on CIFAR-10 while requiring no weight modification, no training overhead, and no FID cost.

2 Related Work

Model Watermarking. Prior machine learning model watermarking methods typically embed an identifiable pattern into model behavior, most notably through trigger-based backdoors or carefully constructed inputs that induce a predetermined response [Adi et al., 2018, Zhang et al., 2018, Fernandez et al., 2023, Wen et al., 2023]. Such approaches provide a direct verification interface when the watermark remains intact, but they generally require explicit intervention during training and can be sensitive to subsequent model editing or "cleaning." *In contrast to embedding an external trigger, our setting treats a private member evidence set as the "watermark" and derives evidence from the model's intrinsic memorization behavior, avoiding modifications that introduce new behavioral artifacts.*

Membership Inference Attacks. Membership inference attacks (MIA) aim to decide whether an input was part of a model's training data [Shokri et al., 2017, Yeom et al., 2018, Salem et al., 2019]. For diffusion models, recent work instantiates MIA using reconstruction-based signals [Carlini et al., 2023, Duan et al., 2023, Kong et al., 2024], loss/likelihood-based criteria [Matsumoto et al., 2023, Hu and Pang, 2023], or scalable quantile-regression formulations [Bertran et al., 2023, Tang et al., 2024]. These methods are primarily evaluated as per-example decision rules (binary classification or scoring) and are typically reported via attack metrics. Ownership verification, however, requires comparative, population-level evidence: the ownership claim must be supported by aggregate separation on a

designated evidence set, calibrated against reference models that represent the null hypothesis, and summarized via statistically interpretable tests rather than single-sample predictions. *Our framework uses MIA-style signals as the underlying measurements, but creates an auditable verification procedure with public references and formal decision criteria.*

Diffusion Model IP Protection. Diffusion-specific IP protection has developed along two complementary lines. *Training-time watermarking* embeds a recoverable signal during training; representatives for unconditional diffusion include WDM [Peng et al., 2023] and Zhao et al. [Zhao et al., 2023], while Stable Diffusion has attracted dedicated designs such as Tree-Ring [Wen et al., 2023] (a Fourier-domain noise pattern), SleeperMark [Wang et al., 2025] (a trigger-prompt-coupled, LoRA-resistant watermark), Stable Signature [Fernandez et al., 2023], RingID [Ci et al., 2024], GaussMarker [Li et al., 2025], and WOUAF [Kim et al., 2024]. *Inference-time dataset inference* is a complementary, training-free line: CDI [Dubiński et al., 2024] aggregates seven membership inference attacks through a meta-classifier to produce a dataset-level membership score. *Two framework-level distinctions separate our work from these prior approaches.* (i) *Training-time vs. inference-time:* the training-time methods above all require a modified training pipeline and incur measurable FID cost, whereas our framework attaches to an already-trained model with $\Delta\mathrm{FID}{=}0$ and zero training overhead. (ii) *Open-ended vs. private-set verification:* CDI shares our inference-time, no-FID-penalty regime but operates against open-ended public-DDPM membership, whereas we leverage a private member evidence set held by the owner to enable calibrated population-level hypothesis testing under controlled false-positive budgets. Empirical comparison against WDM, Zhao, and CDI is given in Section 5.4 (Tables 6 and 5).

3 Problem Setup

3.1 Background and Problem Formulation

Diffusion Models. A diffusion model defines a forward noising process and a learned reverse process. Given a clean sample $x_0 \sim q(x_0)$, the forward process produces a noisy version at timestep $t \in \{1, \ldots, T\}$:

$$x_t = \sqrt{\bar{\alpha}_t}\, x_0 + \sqrt{1 - \bar{\alpha}_t}\, \epsilon, \quad \epsilon \sim \mathcal{N}(0, I), \tag{1}$$

where $\bar{\alpha}_t = \prod_{s=1}^{t} \alpha_s$ is the cumulative noise schedule. A neural network ϵ_θ is trained to recover the added noise by minimizing:

$$\mathcal{L}(\theta) = \mathbb{E}_{x_0,\, \epsilon,\, t}\big[\|\epsilon - \epsilon_\theta(x_t, t)\|^2\big]. \tag{2}$$

At inference, new samples are drawn by iteratively denoising from $x_T \sim \mathcal{N}(0, I)$ using the learned ϵ_θ.

It has been empirically established that diffusion models memorize portions of their training data Carlini et al. [2023], Somepalli et al. [2023]: the reconstruction error on training examples is systematically lower than on unseen samples drawn from the same distribution. This memorization gap persists across architectures and dataset scales, and serves as the foundational signal that our framework exploits for ownership verification.

From Membership Inference to Ownership Verification. Membership inference attacks (MIAs) leverage the memorization gap to determine whether a single query sample belongs to the training set of a target model Duan et al. [2023], Kong et al. [2024]. We state this formally:

Definition 1 (Membership Inference Attack)**.** *Given a target model M and a query sample x, a membership inference attack is a binary classifier $\mathcal{A} : (M, x) \to \{0, 1\}$ that predicts whether $x \in \mathcal{D}_{\text{train}}$.*

However, low reconstruction error on a single sample does not constitute evidence of ownership, as the sample may simply be inherently easy to reconstruct for any model. Ownership verification demands a stronger, population-level claim: the suspect model must exhibit statistically significant memorization on a *designated evidence set* when compared against public baselines that share no training overlap with that set.

Definition 2 (Ownership Verification)**.** *Let $\mathcal{M}_A$ denote a base model trained by the owner, $\mathcal{M}_B$ a suspect model held by the adversary, and $\mathcal{W} = \{x_1, \ldots, x_{|\mathcal{W}|}\}$ a private evidence set held by*

the owner such that $\mathcal{W} \subset \mathcal{D}_{\text{train}}(\mathcal{M}_A)$. An ownership verification protocol is a hypothesis test $\mathcal{V} : (\mathcal{M}_B, \mathcal{W}, \{M_{ref}^{(i)}\}) \to \{\texttt{owned}, \texttt{independent}\}$ that determines whether $\mathcal{M}_B$ is derived from $\mathcal{M}_A$ by comparing the reconstruction behavior of $\mathcal{M}_B$ on $\mathcal{W}$ against a set of public baseline models $\{M_{ref}^{(i)}\}_{i=1}^{|\mathcal{R}|}$ at a controlled false-positive rate α.

This formulation shifts the task from sample-level binary classification to a calibrated, population-level statistical test, which is the protocol we develop in the following sections.

3.2 Threat Model

We consider three principals in the ownership verification setting. The **owner** O trains a model $\mathcal{M}_A$ on $\mathcal{D}_{\text{train}}$, which contains a private watermark set $\mathcal{W} \subset \mathcal{D}_{\text{train}}$ known only to O. The **adversary** obtains $\mathcal{M}_A$ through theft or leakage, and produces a derivative model $\mathcal{M}_B$ by fine-tuning; $\mathcal{M}_B$ is the suspect model submitted for verification. The **verifier** V, given $\mathcal{M}_A$, $\mathcal{M}_B$, a set of public baseline models, and the private set $\mathcal{W}$, determines whether the ownership claim is valid.

We assume the adversary has white-box access to $\mathcal{M}_A$ and may possess in-domain data disjoint from $\mathcal{W}$. The adversary is aware that a verification protocol exists but cannot identify $\mathcal{W}$: even if $\mathcal{D}_{\text{train}}$ were leaked, selecting the exact $|\mathcal{W}|$-subset from N training samples is feasible with probability at most

$$P(\text{exact match}) = \frac{1}{\binom{N}{|\mathcal{W}|}} \;\leq\; \left(\frac{|\mathcal{W}|}{N}\right)^{|\mathcal{W}|}, \tag{3}$$

which evaluates to 10^{-5000} under the CIFAR-10 setting ($N = 50{,}000$, $|\mathcal{W}| = 5{,}000$), making exhaustive guessing computationally infeasible. This is a counting argument and does not claim resistance to side-channel leakage of $\mathcal{W}$. The adversary is permitted to apply fine-tuning, distillation, and pruning to $\mathcal{M}_A$; full retraining from scratch falls outside our scope, as it severs provenance entirely.

Upon receiving a dispute, the verifier computes t-error scores on $\mathcal{W}$ for both the suspect model and public baselines, then evaluates the two-point criteria (Section 4.4), outputting VERIFIED or REJECTED together with the supporting statistical evidence (p-values and Cohen's d). The verifier is assumed to have white-box access to both models, consistent with standard practice in IP litigation, regulatory audits, and contractual audit clauses.

4 Method

The verification framework has four components: (1) a *t-error* score measuring reconstruction quality via forward-reverse diffusion, (2) multi-timestep aggregation for stable per-sample scores, (3) Gaussian quantile regression for FPR-controlled thresholds, and (4) a two-point statistical protocol comparing the suspect model against public baselines.

4.1 t-error Scoring

Diffusion models reconstruct their training data with lower error than unseen data [Carlini et al., 2023]. We exploit this *memorization gap* by measuring reconstruction quality at a given noise level t.

Specifically, given a clean sample x_0, we first corrupt it via the forward diffusion schedule:

$$x_t = \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{1 - \bar{\alpha}_t} \cdot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I) \tag{4}$$

where x_t is the noised image at timestep $t \in \{1, \ldots, T\}$, ϵ is i.i.d. Gaussian noise, and $\bar{\alpha}_t = \prod_{s=1}^{t} \alpha_s$ is the cumulative noise coefficient under a cosine schedule [Dhariwal and Nichol, 2021]. The model then predicts the added noise $\hat{\epsilon}_\theta(x_t, t)$, from which we recover a single-step reconstruction:

$$\hat{x}_0 = \frac{x_t - \sqrt{1 - \bar{\alpha}_t} \cdot \hat{\epsilon}_\theta(x_t, t)}{\sqrt{\bar{\alpha}_t}} \tag{5}$$

We define the *t-error*, denoted $e_t(x_0)$, as the per-pixel mean squared error between the original and reconstructed images:

$$e_t(x_0) = \frac{\|x_0 - \hat{x}_0\|_2^2}{H \times W \times C} \tag{6}$$

Algorithm 1 t-error Computation

Require: Model $\mathcal{M}$, sample x_0, timesteps $\mathcal{T} = \{t_1, \ldots, t_K\}$, noise schedule $\{\bar{\alpha}_t\}$
Ensure: Aggregated score $s(x_0)$
  errors $\leftarrow$ []
  for $t \in \mathcal{T}$ **do**
    $\epsilon \sim \mathcal{N}(0, I)$
    $x_t \leftarrow \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{1 - \bar{\alpha}_t} \cdot \epsilon$
    $\hat{\epsilon} \leftarrow M_\theta(x_t, t)$
    $\hat{x}_0 \leftarrow (x_t - \sqrt{1 - \bar{\alpha}_t} \cdot \hat{\epsilon}) / \sqrt{\bar{\alpha}_t}$
    $e_t \leftarrow \|x_0 - \hat{x}_0\|_2^2 / (H \times W \times C)$
    errors.append(e_t)
  end for
  return $s(x_0) \leftarrow Q_{0.25}(\texttt{errors})$

where H, W, C denote the image height, width, and number of channels, respectively. The key property underlying our framework is the *memorization effect*: models trained on a particular dataset exhibit systematically lower reconstruction error on their training examples, i.e., $\mathbb{E}_{x \in \mathcal{D}_{\text{train}}}[e_t(x)] < \mathbb{E}_{x \notin \mathcal{D}_{\text{train}}}[e_t(x)]$. This persistent gap, which survives across architectures and dataset scales, provides the foundational signal for our ownership verification protocol.

For latent diffusion models (e.g., Stable Diffusion), the t-error is computed in VAE latent space rather than pixel space; we defer the full treatment to Section 5.5.

4.2 Multi-Timestep Aggregation

A single timestep provides a noisy estimate of reconstruction quality, as the magnitude of $e_t(x)$ varies substantially with the noise level t. To obtain a stable per-sample statistic, we aggregate t-errors over K uniformly sampled timesteps into a single score $s(x)$ via the 25th percentile:

$$s(x) = Q_{0.25}\big(\{e_{t_i}(x)\}_{i=1}^{K}\big) \tag{7}$$

We use K=50 by default. The choice of a lower percentile is deliberate: it captures the timesteps where the memorization gap is most pronounced while remaining robust to outlier timesteps that introduce high but uninformative error. Ablations in Section 5.6 confirm that the 25th percentile yields the strongest separation among common aggregation strategies, and that performance is stable for $K \geq 25$. The full scoring procedure is summarized in Algorithm 1.

4.3 Gaussian Quantile Regression

Rather than relying on fixed thresholds, we model the *conditional score distribution* for non-members and derive decision boundaries in closed form. Concretely, we apply a log transformation to the aggregated score and parameterize the result as a Gaussian with input-dependent mean $\mu(x)$ and standard deviation $\sigma(x)$:

$$y = \log(1 + s(x)) \sim \mathcal{N}\big(\mu(x),\, \sigma^2(x)\big) \tag{8}$$

where $\mu(x)$ and $\sigma(x)$ are jointly predicted by a neural network f_ψ. The network uses a ResNet-18 backbone to extract image features, which are concatenated with a compact statistical summary $\phi(x) \in \mathbb{R}^3$ consisting of the mean, standard deviation, and ℓ_2 norm of the t-error sequence. An MLP head maps the combined representation to $(\mu, \log \sigma)$. These auxiliary statistics provide a direct view of the t-error magnitude that the image features alone cannot capture, stabilizing the quantile predictions.

We train f_ψ by minimizing the Gaussian negative log-likelihood over N auxiliary non-member samples:

$$\mathcal{L}_{\text{NLL}} = \frac{1}{N} \sum_{i=1}^{N} \left[\frac{(y_i - \mu_i)^2}{2\sigma_i^2} + \log \sigma_i \right] \tag{9}$$

where $y_i = \log(1 + s(x_i))$ is the observed log-score and (μ_i, σ_i) are the corresponding network predictions. Because the conditional distribution is Gaussian, the $(1-\alpha)$-quantile at any target

false-positive rate α admits a closed-form expression:

$$\hat{q}_\tau(x) = \mu(x) + \sigma(x) \cdot \Phi^{-1}(\tau), \quad \tau = 1 - \alpha \tag{10}$$

with Φ^{-1} denoting the standard normal inverse CDF. This is a key practical advantage over pinball-loss quantile regression, which requires training a separate model for every desired quantile level; here, a single model supports arbitrary τ without retraining.

To translate the predicted quantile into a membership decision, we define the *margin* $m_\tau(x)$ as the gap between the predicted quantile and the observed log-score:

$$m_\tau(x) = \hat{q}_\tau(x) - \log(1 + s(x)) \tag{11}$$

A positive margin indicates that the sample's reconstruction error falls below what would be expected of a non-member, suggesting likely membership in the training set.

Finally, to reduce prediction variance, we wrap the above procedure in a bagging ensemble of B=50 models, each trained on a bootstrap sample drawn with replacement from 80% of the auxiliary data. At inference, we average the individual quantile predictions into a single ensemble estimate:

$$\hat{q}_\tau^{\text{ens}}(x) = \frac{1}{B} \sum_{b=1}^{B} \hat{q}_\tau^{(b)}(x) \tag{12}$$

where $\hat{q}_\tau^{(b)}(x)$ is the quantile from the b-th ensemble member. The ensemble margin $m_\tau^{\text{ens}}(x) = \hat{q}_\tau^{\text{ens}}(x) - \log(1 + s(x))$ serves as the final per-sample statistic; as before, positive values indicate likely membership. Ensembling is embarrassingly parallel and introduces negligible additional inference cost (Appendix C).

4.4 Two-Point Verification Criteria

Given the owner model $\mathcal{M}_A$, the suspect model $\mathcal{M}_B$, a set of public baselines, and the watermark set $\mathcal{W}$, we require both of the following criteria to be satisfied before declaring ownership verified.

Criterion 1: Consistency. If $\mathcal{M}_B$ is indeed derived from $\mathcal{M}_A$, the two models should exhibit similar reconstruction behavior on $\mathcal{W}$. We operationalize this by requiring that a two-sample t-test on their score distributions yields $p > 0.05$, i.e., we cannot reject the null hypothesis that the two models share the same t-error distribution. Failure here suggests that the suspect model has no provenance relationship with the owner model.

Criterion 2: Separation. The owner model must reconstruct watermark samples significantly better than any public baseline. We require both statistical significance ($p < 10^{-6}$, to control for multiple baseline comparisons) and practical significance ($|d| > 2.0$, corresponding to a "very large" effect size [Cohen, 1988]). This dual requirement guards against claims that achieve a low p-value purely through large sample size without meaningful effect.

Ownership is declared **verified** only when both criteria are jointly satisfied. This conjunction ensures that no single metric can produce a false positive: consistency confirms provenance, while separation establishes that the memorization signal is genuine rather than coincidental. The complete verification protocol is formalized in Algorithm 2.

4.5 Paired-Delta Protocol for Derivative Models

When $\mathcal{M}_B$ is a derivative of $\mathcal{M}_A$ (shared base weights via fine-tuning), the raw t-error gap between owner and reference collapses and Algorithm 2's absolute-score separation criterion no longer applies directly. We replace per-sample scores $s(x)$ with paired differences $\delta(x) = s_{\text{tgt}}(x) - s_{\text{ref}}(x)$ and apply C1/C2 to the δ distribution (see Section 5.5 for full treatment).

5 Experiments

5.1 Experimental Setup

We evaluate our framework on four datasets spanning different resolutions and scales: CIFAR-10 and CIFAR-100 (both 32×32, 50k training images with 5k reserved as the watermark set), STL-10

Algorithm 2 Ownership Verification Protocol

Require: Models M_A, M_B, baselines $\{M_{\text{ref}}^{(i)}\}_{i=1}^{|\mathcal{R}|}$, watermark set $\mathcal{W}$
Ensure: VERIFIED or REJECTED
  $S_A \leftarrow \{s_{M_A}(x) : x \in \mathcal{W}\}$; $S_B \leftarrow \{s_{M_B}(x) : x \in \mathcal{W}\}$
  O1: $p_1 \leftarrow \text{TTest}(S_A, S_B)$; **if** $p_1 < 0.05$ **return** REJECTED
  O2: for each reference $i \in \{1, \ldots, |\mathcal{R}|\}$:
    $S_{\text{ref}}^{(i)} \leftarrow \{s_{M_{\text{ref}}^{(i)}}(x) : x \in \mathcal{W}\}$
    $p_2^{(i)}, d^{(i)} \leftarrow \text{TTest}(S_A, S_{\text{ref}}^{(i)}), \text{CohenD}(S_A, S_{\text{ref}}^{(i)})$
  if $\exists i : p_2^{(i)} > 10^{-6}$ **or** $|d^{(i)}| < 2.0$ **then**
    return REJECTED
  end if
  return VERIFIED

Table 1: Key hyperparameters for reproducibility.

Component	Parameter	Value
DDIM	Iterations	400k
	Batch size / LR	128 / 2×10^{-4}
	EMA decay	0.9999
Gaussian QR	Ensemble B	50
	Bootstrap ratio	0.8
	Epochs / patience	50 / 10
Stable Diffusion	Base model	SD v1.4
	LoRA rank / α	64 / 64, 256 / 256
	Fine-tune steps / LR	10,000 / 1×10^{-4}
	Private set size	1,000 COCO 2014
	Adversary fine-tune (steps / lr)	2,000 / 5×10^{-5}
	Adversary LoRA target / rank	cross-attention / 64 or 256
	Adversary loss	ϵ-prediction MSE
	Adversary training data	1,000 COCO images disjoint from $\mathcal{W}$

(96×96, 5k training images with 1k as watermark), and CelebA (64×64, 162k training images with 5k as watermark). All splits are deterministic and stratified with a fixed seed to ensure reproducibility.

For each dataset, we train a DDIM model using a U-Net backbone with 128 base channels, a cosine noise schedule, and the AdamW optimizer at a learning rate of 2×10^{-4} with EMA decay 0.9999. Training runs for 400k iterations in all cases. The Gaussian QR component uses an ensemble of B=50 models, each trained on a bootstrap sample of the auxiliary non-member data. Table 1 summarizes the key hyperparameters; full architectural and optimization details are deferred to Appendix A.

As public baselines for the verification protocol, we use pretrained checkpoints from HuggingFace: `google/ddpm-cifar10-32` serves as the reference for both CIFAR-10 and CIFAR-100, `google/ddpm-ema-bedroom-256` (resized) for STL-10, and `google/ddpm-celebahq-256` for CelebA. These baselines are trained on disjoint data from our watermark sets and represent the null hypothesis in the two-point verification protocol: a model that has never seen the owner's private evidence should exhibit substantially higher reconstruction error on that set.

We additionally evaluate on Stable Diffusion v1.4 to test whether the framework extends to large-scale parameter-efficient fine-tuning settings. Owner models are obtained by fine-tuning SD v1.4 on 1,000 private COCO 2014 images via LoRA (rank $r \in \{64, 256\}$) or full fine-tuning; full SD experimental setup and the paired-delta verification protocol are described in Section 5.5.

5.2 RQ1: Verification Effectiveness

Can private member evidence sets distinguish owner models from public baselines?

Table 2: Multi-reference ownership verification across 4 datasets. Each owner model is tested against 3 public reference models grouped into semantic-similar DDPM / cross-domain LSUN / random initialization. Verification requires *every* reference to satisfy $|d| > 2.0$ (ALL-must-pass). All 12 (dataset, reference) pairs satisfy the threshold; verification passes on all four datasets. All pairwise $p < 10^{-100}$. †Conservative reference (LSUN-Church, replacing LSUN-Bedroom to address domain-match concerns from prior review).

Dataset	Semantic-similar DDPM $\lvert d\rvert$	Cross-domain LSUN DDPM $\lvert d\rvert$	Random init $\lvert d\rvert$
CIFAR-10	24.1	40.9	63.9
CIFAR-100	18.5	31.6	63.0
STL-10	116.7	29.3†	59.3
CelebA	26.3	32.0	90.0

References per dataset: CIFAR-10/100 use `google/ddpm-cifar10-32`, `google/ddpm-bedroom-256`, random-32. STL-10 uses `google/ddpm-cifar10-32`, `google/ddpm-church-256`†, random-96. CelebA uses `google/ddpm-celebahq-256`, `google/ddpm-bedroom-256`, random-64.

Table 3: Membership inference performance. DDIM rows use Gaussian QR margins (TPR computed at fixed FPR with 95% bootstrap CI, 200 iterations). SD rows use latent-space paired-delta scores evaluated on 1,000 members + 1,000 non-members; TPR at 0.01% FPR is omitted as it falls below the discretization floor (1 false positive at this split size).

Dataset	AUC	TPR@0.1%	TPR@0.01%
CIFAR-10	0.987 ± 0.002	0.89 ± 0.03	0.71 ± 0.05
CIFAR-100	0.982 ± 0.003	0.85 ± 0.04	0.65 ± 0.06
STL-10	0.968 ± 0.005	0.78 ± 0.05	0.55 ± 0.07
CelebA	0.975 ± 0.004	0.82 ± 0.04	0.60 ± 0.06
SD v1.4 (LoRA $r=64$)	0.995	0.582	—
SD v1.4 (LoRA $r=256$)	0.999	0.879	—
SD v1.4 (Full FT)	0.999	0.847	—

Table 2 reports Cohen's d against three categories of public reference models (semantic-similar DDPM, cross-domain LSUN DDPM, and random initialization) under the ALL-must-pass criterion.

Results: Table 2 reports verification under the ALL-must-pass criterion: every public reference model must individually satisfy $|d| > 2.0$. We test each owner model against 3 reference models spanning a semantic-similar DDPM, a cross-domain LSUN DDPM, and a random initialization. The minimum $|d|$ across all references ranges from 18.5 (CIFAR-100, weakest dataset) to 29.3 (STL-10 with the LSUN-Church reference), exceeding the threshold by an order of magnitude. For STL-10, we use LSUN-Church as a conservative replacement for LSUN-Bedroom, which a prior reviewer flagged as a potential domain-match confound; the church reference yields $|d| = 29.3$, still $14\times$ above the $|d| > 2.0$ threshold. These margins confirm that the verification signal is an intrinsic property of the training data, not an artifact of a particular reference model choice.

Consistency (C1). Owner and Model B produce statistically indistinguishable mean t-error on $\mathcal{W}$ across all four datasets: CIFAR-10: 28.73 vs. 28.74 ($p = 0.98$); CIFAR-100: 33.35 vs. 33.34 ($p = 0.98$); STL-10: 70.92 vs. 71.38 ($p = 0.71$); CelebA: 63.79 vs. 63.80 ($p = 0.97$). All p-values exceed the 0.05 acceptance threshold, confirming that Model B retains the memorization signature of its owner.

Table 3 reports per-sample membership inference accuracy using Gaussian QR margins.

Threshold Sensitivity: The threshold $|d| > 2.0$ is conservative. Table 4 reports the minimum $|d|$ each dataset achieves and the strictest integer threshold it still passes. The smallest margin (CIFAR-100 at $9.3\times$) leaves substantial headroom above the recommended threshold, confirming that results are not artifacts of threshold tuning.

Verification effectiveness on Stable Diffusion. Verification effectiveness extends beyond the DDIM-from-scratch setting to Stable Diffusion v1.4 with parameter-efficient fine-tuning. Under the paired-delta protocol of Section 5.5, all three owner configurations—LoRA r=64, full fine-tuning,

Table 4: Verification margin per dataset. Each row reports the minimum $|d|$ across all three reference models (semantic-similar / cross-domain LSUN / random initialization), the strictest integer $|d|$ threshold that the dataset still passes, and the multiplicative margin above our recommended threshold $|d| > 2.0$.

Dataset	Min $\lvert d\rvert$ achieved	Strictest threshold passed	Margin above $\lvert d\rvert > 2.0$
CIFAR-10	24.1	$\lvert d\rvert > 24$	12.0×
CIFAR-100	18.5	$\lvert d\rvert > 18$	9.3×
STL-10	29.3	$\lvert d\rvert > 29$	14.7×
CelebA	26.3	$\lvert d\rvert > 26$	13.2×

and LoRA r=256—satisfy both C1 (consistency, $p > 0.05$ on every adversary) and C2 (separation, $|d| \in \{2.38, 2.43, 2.55\}$) across domain-shift and task-shift adversaries (Table 7), producing a ✓ verdict on all six (owner, adversary) pairs.

5.3 RQ2: Post-Theft Robustness Across Multiple Attacks and Methods

A practical adversary who obtains a stolen model will typically fine-tune it to a new domain or task before deployment. The ownership signal must survive such adaptation. To evaluate this, we simulate post-theft modification via MMD fine-tuning [Gretton et al., 2012]. MMD fine-tuning contrasts CLIP features of the model's DDIM-generated samples against those of real training images sampled fresh each step from $\mathcal{D}_{\text{train}}$ (which includes $\mathcal{W}$), via a cubic polynomial kernel. Starting from Model A, we fine-tune for 500 iterations using 10-step DDIM sampling, CLIP ViT-B/32 as the feature encoder, a learning rate of 5×10^{-6}, and EMA decay 0.9999.

The resulting Model B retains nearly identical t-error to Model A on the watermark set (Table 2), and both verification criteria remain satisfied across every dataset. This robustness is not coincidental: because the loss operates at the distribution level rather than imposing a per-image reconstruction objective on $\mathcal{W}$, the memorization gap on individual $\mathcal{W}$ examples persists. While WDM and Zhao also withstand these attacks in our evaluation (Table 5), MiO achieves this robustness without the training-time cost and FID degradation that those methods incur (Table 6). We note, however, that an adversary who retrains from scratch or specifically targets memorization erasure falls outside our threat model.

Table 5 stress-tests the three diffusion-model baselines under four post-theft attacks: MMD fine-tuning and standard SGD fine-tuning (both 500 iterations at lr$=5\times10^{-6}$ on clean CIFAR-10), and additive Gaussian weight noise at $\sigma=10^{-3}$ and $\sigma=10^{-2}$. All three methods pass every measured attack cell at our chosen strengths; MiO's distinguishing value therefore lies not in exclusive robustness but in delivering comparable robustness with no training overhead, no FID cost, and no modification to the target model's weights (Table 6). CDI's robustness under fine-tuning is not reported in the original paper and was not re-run on adversary checkpoints here; we therefore omit CDI from Table 5 and defer that evaluation to future work.

SD adversary robustness. Post-theft robustness extends to the SD-PEFT setting under the paired-delta protocol. Across both domain-shift fine-tuning on disjoint COCO and task-shift fine-tuning on synthetic data, the verdict holds uniformly for all three owner configurations (Table 7). Under the more stringent 2,000-step LoRA fine-tune on disjoint COCO data (rank $r \in \{64, 256\}$, lr 5×10^{-5}, cross-attention modules only, ϵ-prediction MSE) that simulates post-theft adaptation, MiO retains the verification verdict for both r=64 and r=256 adapters, matching SleeperMark's trigger-based bit-accuracy threshold on the same threat model (Table 8).

5.4 Controlled Baseline Comparison

Whereas Section 5.3 shows that MiO matches WDM and Zhao under post-theft attacks, this section focuses on the orthogonal question of capability cost: training overhead, FID degradation, and weight modification. We compare MiO against three existing methods on CIFAR-10: two training-time watermarking methods (WDM [Peng et al., 2023], which learns a secondary watermark diffusion process alongside the standard DDPM objective, and Zhao et al. [Zhao et al., 2023], which embeds a StegaStamp bit-string into training images before fitting an EDM model) and one inference-time

Table 5: Robustness of watermark and ownership verification under four post-theft attacks on CIFAR-10. ✓ indicates the method passes its native pass/fail threshold; "—" indicates not applicable. Pass thresholds: pixel accuracy > 0.9 for WDM (on the WDP-extracted watermark image); bit accuracy > 0.9 for Zhao (on $N = 1000$ Heun-18 generated samples decoded with the ground-truth 64-bit StegaStamp fingerprint); two-point verification (C1: $p > 0.05$, C2: $|d| > 2.0$) for MiO. WDM/Zhao thresholds are MiO-defined; the original papers do not specify robustness thresholds. Attack configurations: MMD-FT and SGD-FT each run for 500 iterations at $\text{lr} = 5 \times 10^{-6}$ on clean CIFAR-10. SGD-FT uses each method's native training objective (ε-prediction MSE for the DDPM-parameterized WDM and MiO; σ-weighted $\hat{x}_0$-prediction MSE for Zhao's EDM backbone) to avoid introducing a parameterization mismatch. Gaussian weight noise is applied in-place to all parameter tensors of each method's native network. ‡Zhao × MMD-FT requires a Heun-based differentiable sampler not implemented in our DDIM-based MMD fine-tuning loop; deferred to future work.

Method	MMD-FT	SGD-FT	$\sigma = 10^{-3}$	$\sigma = 10^{-2}$
WDM	✓	✓	✓	✓
Zhao†	—‡	✓	✓	✓
MiO	✓	✓	✓	✓

†Uses EDM backbone (not DDPM/DDIM).

MIA-aggregation method (CDI [Dubiński et al., 2024]), which aggregates seven complementary membership inference attacks (denoising loss, SecMI-stat, PIA, PIAN, gradient masking, noise optimization, multiple-loss) through a logistic-regression meta-classifier. All three baselines have publicly available official implementations. WDM and Zhao are retrained on the same 50K CIFAR-10 images so that any differences reflect the methods themselves rather than data or scale effects; CDI is evaluated on the public `google/ddpm-cifar10-32` checkpoint using the authors' released codebase [Dubiński et al., 2024] with three trivial patches required for the aggregation stage to execute end-to-end (Appendix A.1); 5K members + 5K non-members under the default 7-attack configuration. We exclude DeepMarks [Chen et al., 2019], whose GMM-regularised weight fingerprinting was designed for discriminative classifiers and has no reference implementation for generative architectures.

Table 6 summarises the four methods along four capability dimensions: whether the method modifies the target model's weights, the additional training cost it incurs, its FID cost (the generation-quality penalty attributable to the method), and its per-sample membership AUC. WDM and Zhao both require a full (re)training pass with a modified objective and introduce a measurable FID cost (+7.21 and +3.07 respectively). CDI and MiO are both inference-only: they leave the target model's weights unchanged, so their FID cost is zero by construction and they incur no training overhead.

On the one dimension where CDI and MiO are directly comparable— per-sample membership AUC on CIFAR-10—MiO achieves 0.987 against CDI's 0.586, a substantial gap even after accounting for the fact that CDI and MiO evaluate on different base checkpoints (`google/ddpm-cifar10-32` for CDI; the owner's custom DDIM for MiO). The gap reflects the structural advantage of operating against a private member evidence set held by the owner: when the defender controls a known member subset, per-sample detection becomes tractable at stringent false-positive rates, whereas open-ended dataset inference operates in a harder setting.

5.5 Extension to Stable Diffusion

The preceding experiments use DDIM models trained from scratch. Modern practice, however, centres on adapting large pretrained models via parameter-efficient fine-tuning. We therefore evaluate MiO on Stable Diffusion v1.4, where the owner fine-tunes the base model on a private set of 1,000 COCO 2014 images using LoRA or full fine-tuning, and the adversary produces derivative models via continued LoRA training on disjoint data.

Measurement recipe. Because SD operates in VAE latent space, the t-error can be measured either in latent space (before the decoder) or in pixel space (after decoding). We find that *latent-space scoring with per-image caption conditioning* yields $\text{AUC} = 0.9956$ on the owner model, compared to 0.959 for pixel-space scoring—a 3.7 percentage-point gap. Two factors explain this difference. First, the VAE decode step introduces a lossy reconstruction floor that dilutes the membership signal in pixel

Table 6: Capability matrix of diffusion-model IP protection methods on CIFAR-10. ✓/✗ indicate whether the method has the property; — indicates the column is not applicable to the method. FID cost values for WDM and Zhao are reported under their original sampling protocols (WDM: 50-step DDIM; Zhao: 18-step Heun on EDM backbone). CDI and MiO are inference-only and do not alter model weights; their FID cost is 0 by construction. Per-sample AUC is reported on CIFAR-10: CDI evaluated on `google/ddpm-cifar10-32` (5K members + 5K non-members); MiO evaluated on its owner DDIM with a 5K private watermark set.

Method	Modifies weights?	Training overhead	FID cost	Per-sample AUC
WDM	✓	$\sim 1\times$	$+7.21$	—
Zhao†	✓	$\sim 1\times$	$+3.07$	—
CDI	✗	0	0	0.586
MiO (Ours)	✗	**0**	**0**	**0.987**

†Uses EDM backbone (not DDPM/DDIM).

Table 7: Ownership verification on SD v1.4 (1,000 private COCO images, paired-delta protocol). C1: consistency ($p > 0.05$). C2: separation ($|d| > 2.0$). Adversaries: domain-shift fine-tuning on disjoint COCO and task-shift on synthetic data. Per-adversary results via paired-delta protocol. $|d|_{\text{mem}}$: effect size between the owner's score distributions on member and non-member samples in $\mathcal{W}$, via standard absolute Cohen's d.

Owner	Adversary	C1	C2 ($\lvert d\rvert$)	AUC	$\lvert d\rvert_{\text{mem}}$	Verdict
LoRA $r=64$	Domain-shift	✓ (p=0.73)	✓ (2.43)	0.986	2.83	✓
LoRA $r=64$	Task-shift	✓ (p=0.34)	✓ (2.43)	0.980	2.70	✓
Full FT	Domain-shift	✓ (p=0.71)	✓ (2.38)	0.997	3.22	✓
Full FT	Task-shift	✓ (p=0.18)	✓ (2.38)	0.995	3.13	✓
LoRA $r=256$	Domain-shift	✓ (p=0.54)	✓ (2.55)	0.998	3.51	✓
LoRA $r=256$	Task-shift	✓ (p=0.22)	✓ (2.55)	0.997	3.46	✓

space. Second, SD learned image-caption associations during training; using empty prompts bypasses the cross-attention pathway where LoRA adapters operate, reducing the observable memorization gap. All SD results below use the latent-plus-caption configuration.

Verification results. Table 7 reports ownership verification across three owner configurations. Because derivative models share the base SD weights, the owner and reference produce nearly identical raw t-errors, and Algorithm 2's absolute-score separation criterion does not reach its threshold directly. We therefore adapt the protocol to use *paired differences*: for each image $x \in \mathcal{W}$, we compute $\delta(x) = s_{\text{tgt}}(x) - s_{\text{ref}}(x)$ and apply the two-point criteria to the δ distribution, which cancels per-image variance shared between models.

Owner-level global membership detection. For reference, owner-level membership detection on the watermark set $\mathcal{W}$ (1,000 members vs. 1,000 non-members, global evaluation aggregated across both adversary fine-tuning directions) achieves AUC $= 0.995/0.999/0.999$ and $|d| = 3.26/3.93/3.41$ for LoRA r$=$ 64 / LoRA r$=$ 256 / Full FT respectively. Per-adversary verification results are aggregated using the paired-delta protocol, which is invariant to adversary fine-tuning direction.

All three owner configurations (LoRA r$=$ 64, Full FT, LoRA r$=$ 256) satisfy both criteria (C1 consistency and C2 separation) under the paired-delta verification protocol, producing a ✓ verdict on every adversary type considered.

5.6 RQ3: Score Stability and Threshold Calibration

Several design choices underlie the verification pipeline—how to collapse a sequence of per-timestep errors into a single score, how many timesteps are needed for that score to be reliable, and whether the Gaussian assumption in our quantile regression is worth the modelling commitment. We address each in turn on CIFAR-10; ensemble size and wall-clock breakdowns appear in Appendix B.

Table 8: Watermark robustness on SD v1.4: SleeperMark vs. MiO. **Threat**: adversary obtains the watermarked base model and applies 2,000-step LoRA fine-tuning (lr 5×10^{-5}, cross-attention modules, ϵ-prediction MSE on 1,000 disjoint COCO images). **Success criterion**: SleeperMark passes if bit-accuracy > 0.9; MiO passes if two-point criteria (C1 + C2) are satisfied. Both methods pass their respective thresholds under LoRA adaptation.

Method	LoRA $r=64$	LoRA $r=256$
SleeperMark	✓	✓
MiO (Ours)	✓	✓

Table 9: Aggregation ablation on CIFAR-10. Lower percentiles sharpen separation up to a point; Q25 offers the best trade-off before outlier sensitivity sets in. Ablation conducted on CIFAR-10; SD results use Q25 and are consistent with this ranking (see Section 5.5).

Aggregation	$\lvert d\rvert$
Mean	18.2
Q10	21.5
Q25 (ours)	**23.93**
Median	19.8

Table 9 compares four ways of summarising the t-error sequence. The mean, which weights every timestep equally, washes out the signal at noise levels where the memorisation gap is widest; the median is slightly more selective but still averages over a broad middle band. Moving to lower percentiles helps because the ownership signal concentrates at moderate noise levels—timesteps at which the model reconstructs its own training data with noticeably lower error than unseen samples. Q25 captures this sweet spot, reaching $|d| = 23.93$. Pushing further to Q10 begins to pick up outlier timesteps whose high variance hurts more than their selectivity helps, and separation drops back slightly.

The number of sampled timesteps K matters less than one might expect. Varying K from 10 to 100, we find that $|d|$ plateaus once $K \geq 25$ and changes negligibly beyond $K = 50$. We therefore fix $K = 50$ throughout; practitioners under tighter compute budgets can halve this to $K = 25$ with only marginal loss.

Finally, we ask whether the Gaussian distributional assumption in our quantile regression buys anything over a distribution-free alternative. Pinball-loss QR optimises directly for a target quantile τ without assuming a parametric form, but it requires a separate model for every τ of interest. On CIFAR-10, at $\tau \in \{0.999,\ 0.9999\}$ (FPR $= 0.1\%$ and 0.01%), the two approaches yield comparable TPR. The practical difference is flexibility: Gaussian QR trains once and evaluates any τ in closed form, cutting both the parameter count and the retraining burden roughly in half. For a verifier who may need to assess ownership claims at several significance levels—exactly the setting envisioned by our protocol—this is a meaningful convenience.

6 Discussion and Limitations

We frame scope conditions rather than limitations.

White-box Access: Our method applies where model weights are available, standard in IP litigation (courts compel disclosure), regulatory audits (GDPR/AI Act), partnership disputes, and increasingly common open-weight models (LLaMA, Mistral).

Watermark Retention: Verification requires the owner to retain the watermark set, analogous to keeping receipts for warranty claims.

Retroactive Claims: Our protocol assumes the owner's evidence set $\mathcal{W}$ is established through normal training-data curation prior to any dispute, not selected post-hoc from a released training corpus. The combinatorial argument in Section 3 bounds the probability that an adversary observing $\mathcal{D}_{\text{train}}$ ex post can guess the exact $|\mathcal{W}|$-subset (e.g., 10^{-5000} for the CIFAR-10 setting), but this bound is meaningful only when $\mathcal{W}$ predates the dispute and the owner can demonstrate its prior existence

through normal-course evidence (training logs, dataset versioning, internal documentation). The protocol does not by itself adjudicate the temporal authenticity of $\mathcal{W}$; that determination is delegated to the same evidentiary processes that govern other intellectual-property disputes.

Baseline Availability: Results are strongest with domain-matched public baselines; for novel domains, general-purpose baselines provide conservative bounds.

Gaussian Assumption: The Gaussian model provides a principled approximation; empirical calibration shows quantile estimates remain accurate under moderate misspecification.

Distillation Boundary: Our threat model assumes the adversary retains the original training distribution via weight-preserving operations (fine-tuning, pruning, quantisation). Knowledge distillation—generating synthetic data from the stolen model and training a fresh student—violates this assumption and verification correctly reports no ownership (FAIL via C1 consistency rejection; Appendix E). This is a principled boundary shared by proactive watermarking methods: WDM and Zhao embed signals into model weights, which distillation with fresh initialisation erases equally.

What This Work Does NOT Address: API-only verification (we focus on white-box), full retraining adversaries (our threat model assumes fine-tuning), and adversaries specifically optimizing to remove memorization signals.

7 Conclusion

We presented a novel ownership verification framework for diffusion models using t-error based reconstruction disparity and Gaussian Quantile Regression. By formalizing the distinction between membership inference and ownership verification, we established that public baselines and statistical hypothesis testing are *necessary* components for ownership claims. Our two-point verification criteria (consistency, separation) provide rigorous, reproducible verification. Experiments across four datasets demonstrate $|d| > 18$ separation against every reference model, with criteria satisfied even after adversarial MMD fine-tuning.

References

Yossi Adi, Carsten Baum, Moustapha Cissé, Benny Pinkas, and Joseph Keshet. Turning your weakness into a strength: Watermarking deep neural networks by backdooring. In *USENIX Security Symposium*, pages 1615–1631, 2018.

Martin Bertran, Shuai Tang, Michael Kearns, Jamie Morgenstern, Aaron Roth, and Zhiwei Steven Wu. Scalable membership inference attacks via quantile regression. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023.

Nicholas Carlini, Jamie Hayes, Milad Nasr, Matthew Jagielski, Vikash Sehwag, Florian Tramèr, Borja Balle, Daphne Ippolito, and Eric Wallace. Extracting training data from diffusion models. In *Proceedings of the 32nd USENIX Security Symposium (USENIX Security 2023)*, pages 5253–5270, 2023.

Huili Chen, Bita Darvish Rohani, and Farinaz Koushanfar. DeepMarks: A secure fingerprinting framework for digital rights management of deep learning models. In *ACM International Conference on Multimedia Retrieval (ICMR)*, pages 105–113, 2019.

Hai Ci, Pei Yang, Yiren Song, and Mike Zheng Shou. Ringid: Rethinking tree-ring watermarking for enhanced multi-key identification. In *Computer Vision – ECCV 2024*, volume 15086 of *Lecture Notes in Computer Science*, pages 338–354, 2024.

Jacob Cohen. *Statistical Power Analysis for the Behavioral Sciences*. Lawrence Erlbaum Associates, 2nd edition, 1988.

Prafulla Dhariwal and Alexander Nichol. Diffusion models beat GANs on image synthesis. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 8780–8794, 2021.

Jiaxin Duan, Fanbo Kong, Ruoxi Wang, Karan Gu, et al. Diffusion model membership inference via reverse process discrepancy. In *International Conference on Machine Learning (ICML)*, 2023.

Jan Dubiński, Antoni Kowalczuk, Franziska Boenisch, and Adam Dziedzic. CDI: Copyrighted data identification in diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.

Pierre Fernandez, Guillaume Couairon, Hervé Jégou, Matthijs Douze, and Teddy Furon. The stable signature: Rooting watermarks in latent diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.

Arthur Gretton, Karsten M Borgwardt, Malte J Rasch, Bernhard Schölkopf, Alexander Smola, et al. A kernel two-sample test. *Journal of Machine Learning Research (JMLR)*, 13(1):723–773, 2012.

Xiangming Gu, Chao Du, Tianyu Pang, Chongxuan Li, Min Lin, and Ye Wang. On memorization in diffusion models. *Transactions on Machine Learning Research (TMLR)*, 2025.

Hailong Hu and Jun Pang. Loss and likelihood based membership inference of diffusion models. In *Proceedings of the 26th Information Security Conference (ISC)*, volume 14411 of *Lecture Notes in Computer Science*, pages 121–141. Springer, 2023.

Changhoon Kim, Kyle Min, Maitreya Patel, Sheng Cheng, and Yezhou Yang. WOUAF: Weight modulation for user attribution and fingerprinting in text-to-image diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8974–8983, 2024.

Fei Kong, Jinhao Duan, Ruipeng Ma, Heng Tao Shen, Xiaoshuang Shi, Xiaofeng Zhu, and Kaidi Xu. An efficient membership inference attack for the diffusion model by proximal initialization. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024.

Kecen Li, Zhicong Huang, Xinwen Hou, and Cheng Hong. Gaussmarker: Robust dual-domain watermark for diffusion models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025.

Yugeng Liu, Zheng Li, Michael Backes, Yun Shen, and Yang Zhang. Watermarking diffusion model. *arXiv preprint arXiv:2305.12502*, 2023.

Tomoya Matsumoto, Takayuki Miura, and Naoto Yanai. Membership inference attacks against diffusion models. In *IEEE Security and Privacy Workshops (SPW)*, pages 77–83, 2023.

Sen Peng, Yufei Nie, Yongzhen Zhang, and Jie Xiong. Protecting the intellectual property of diffusion models by the watermark diffusion process. *arXiv preprint arXiv:2306.03436*, 2023.

Ahmed Salem, Yang Zhang, Mathias Humbert, Pascal Berrang, Mario Fritz, and Michael Backes. ML-Leaks: Model and data independent membership inference attacks and defenses on machine learning models. In *Network and Distributed System Security Symposium (NDSS)*, 2019.

Reza Shokri, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. Membership inference attacks against machine learning models. In *IEEE Symposium on Security and Privacy (S&P)*, pages 3–18, 2017.

Gowthami Somepalli, Vasu Singla, Micah Goldblum, Jonas Geiping, and Tom Goldstein. Diffusion art or digital forgery? investigating data replication in diffusion models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6048–6058, 2023.

Shuai Tang, Zhiwei Steven Wu, Sergül Aydöre, Michael Kearns, and Aaron Roth. Membership inference attacks on diffusion models via quantile regression. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, 2024.

Zilan Wang, Junfeng Guo, Jiacheng Zhu, Yiming Li, Heng Huang, and Muhao Chen. Sleepermark: Towards robust watermark against fine-tuning text-to-image diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.

Yuxin Wen, John Kirchenbauer, Jonas Geiping, and Tom Goldstein. Tree-ring watermarks: Fingerprints for diffusion images that are invisible and robust. *arXiv preprint arXiv:2305.20030*, 2023.

Samuel Yeom, Irene Giacomelli, Matt Fredrikson, and Somesh Jha. Privacy risk in machine learning: Analyzing the connection to overfitting. In *IEEE Computer Security Foundations Symposium (CSF)*, pages 268–282, 2018.

Jialong Zhang, Zhongshu Gu, Jiyong Jang, Hui Wu, Marc Ph Stoecklin, Heqing Huang, and Ian Molloy. Protecting intellectual property of deep neural networks with watermarking. In *ACM Asia Conference on Computer and Communications Security (ASIACCS)*, pages 159–172, 2018.

Yunqing Zhao, Tianyu Pang, Chao Du, Xiao Yang, Ngai-Man Cheung, and Min Lin. A recipe for watermarking diffusion models. *arXiv preprint arXiv:2303.10137*, 2023.

A Additional Experimental Details

A.1 Hyperparameters

A.1.1 DDIM Owner Model Training

Table 10: Training hyperparameters for DDIM models.

Parameter	Value
Total iterations	400,000
Batch size	128
Learning rate	2×10^{-4}
Optimizer	AdamW
Weight decay	0.0
EMA decay	0.9999
Noise schedule	Cosine
Timesteps T	1000
U-Net Architecture	
Base channels	128
Channel multipliers	[1, 2, 2, 2]
Attention resolutions	[16]
Dropout	0.1

A.1.2 Gaussian QR Training

Table 11: Gaussian QR training hyperparameters.

Parameter	Value
Backbone	ResNet-18 (pretrained)
Ensemble size B	50
Bootstrap ratio	0.8
Epochs per model	50
Batch size	256
Learning rate	0.001
LR schedule	Cosine annealing
Early stopping patience	10 epochs

A.1.3 Stable Diffusion Owner Fine-tuning

All three configurations share the same 1,000-image private subset of COCO 2014 as the watermark set $\mathcal{W}$; no per-configuration data resampling.

We fine-tune Stable Diffusion v1.4 from the official base checkpoint. Common settings: AdamW optimizer, mixed-precision fp16, batch size 4, gradient accumulation 1, 512×512 resize with center

crop, per-image COCO captions used as text conditioning, cosine learning-rate schedule with 500 warmup steps, seed 42.

Variants:

- **LoRA** r=64 **(A6):** 20,000 steps, lr 1×10^{-4}, rank/$\alpha = 64/64$, target modules {`to_q`, `to_k`, `to_v`, `to_out.0`}, with random horizontal flip during data loading.
- **LoRA** r=256 **(A8):** 20,000 steps, lr 1×10^{-4}, rank/$\alpha = 256/256$, same target modules as A6.
- **Full FT (A7):** 10,000 steps, lr 5×10^{-6}, all 860M U-Net parameters trainable.

A.1.4 Adversary Fine-tuning

The adversary applies LoRA fine-tuning to either the owner's fine-tuned SD model (Table 7) or the SleeperMark watermarked checkpoint (Table 8). Identical settings in both cases: 2,000 steps, AdamW (weight decay 10^{-2}), lr 5×10^{-5}, batch size 4, gradient clipping at norm 1.0, ϵ-prediction MSE loss on 1,000 disjoint COCO images, seed 42. LoRA targets cross-attention modules {`to_q`, `to_k`, `to_v`, `to_out.0`} with rank/$\alpha \in \{64/64,\ 256/256\}$.

A.1.5 WDM Baseline

Architecture: U-Net with 128 base channels, multipliers $[1, 2, 2, 2]$, 3 residual blocks per stage, attention at resolutions $\{16, 8\}$, 4 attention heads, GroupNorm(32), SiLU activation, dropout 0.1. Noise schedule: linear β over T=1000 (note: differs from MiO's cosine schedule). Prediction: ϵ-prediction. Training: 300,000 steps, batch 32, lr 1×10^{-4}, EMA 0.999, AdamW. Watermark-specific: $\gamma_1 = 0.8$ (training blend factor), $\gamma_2 = 0.1$ (training loss weight on watermarked samples; $\gamma_2 = 0.2$ at extraction), watermark seed 1998, fixed trigger image. Loss: $\mathcal{L} = \mathrm{MSE}(\epsilon) + \gamma_2 \cdot \mathrm{MSE}(\epsilon_{\mathrm{wm}})$. FID sampling uses DDIM 50 steps.

A.1.6 Zhao Baseline

Three-stage pipeline. **Stage 1**: train StegaStamp encoder/decoder at 32×32 resolution, batch 64, 100 epochs, 64-bit fingerprint. **Stage 2**: embed a fixed 64-bit fingerprint (seed 42) into all 50K CIFAR-10 training images. **Stage 3**: train EDM with SongUNet (ddpmpp) backbone, 128 channels, multipliers $[2, 2, 2]$, 4 blocks per stage, attention at resolution 16; EDMPrecond head with σ-weighted $\hat{x}_0$-prediction MSE loss; total batch 512, dropout 0.13, augmentation probability 0.12, lr 10^{-3}, EMA 0.5, seed 42. Training duration was early-stopped at 15 Mimg (the original paper trains to 200 Mimg; full duration would require approximately 57 single-GPU-days). The 15 Mimg checkpoint achieves 100% fingerprint bit accuracy on watermark detection, sufficient for the baseline comparison. FID sampling uses Heun 18 steps.

A.1.7 SleeperMark Baseline

We use the pretrained SleeperMark U-Net checkpoint released by the authors (approximately 3.3 GB). No SleeperMark training is performed in our pipeline; the Table 8 results measure robustness of the released checkpoint under the adversary protocol described above.

A.1.8 CDI Reproduction

We re-evaluated CDI [Dubiński et al., 2024] on `google/ddpm-cifar10-32` using the authors' codebase[1] on a single GPU (Python 3.8.18, PyTorch 2.4.1+cu121, diffusers 0.36.0; total wall-clock ≈49 minutes). **Configuration**: 5,000 members and 5,000 non-members sampled from CIFAR-10; batch size 100; seed 0; three-stage Hydra pipeline (features_extraction → scores_computation → evaluation). **Aggregation**: logistic regression (max_iter 1000) over the union of four score-based attacks (denoising loss, SecMI-stat, PIA, PIAN) and three feature-based attacks (gradient masking, noise optimization, multiple-loss). Individual attack AUCs ranged from 0.491 (PIAN) to 0.557 (noise optimization); the aggregated CDI AUC reported in Table 6 is 0.586.

[1] `https://github.com/sprintml/copyrighted_data_identification`

Patches. Three trivial patches were required for the upstream `scores_computation/cdi.py` aggregation stage to execute end-to-end: a missing `import numpy as np`; a `train_dataset["data"]` → `train_dataset.data` fix for the `MIDataset` class API; and a 3D→2D reshape for multi-measurement features. None affects the aggregation algorithm itself.

A.2 Post-theft Attack Configurations

Table 12: MMD fine-tuning hyperparameters.

Parameter	Value
Iterations	500
Learning rate	5×10^{-6}
Batch size	128
DDIM steps	10
t_{start}	900
EMA decay	0.9999
Kernel	Cubic polynomial
Feature extractor	CLIP ViT-B/32 (frozen)

SGD fine-tuning. 500 iterations, lr 5×10^{-6}, batch 128, applied on clean CIFAR-10 using each method's native training objective: ϵ-prediction MSE for the DDPM-parameterized WDM and MiO; σ-weighted $\hat{x}_0$-prediction MSE for Zhao's EDM backbone. The choice of native objectives avoids introducing a parameterization mismatch.

Gaussian weight noise. Additive isotropic Gaussian noise of standard deviation $\sigma \in \{10^{-3}, 10^{-2}\}$ applied in-place to all parameter tensors of each method's native network. No retraining; the perturbed checkpoint is evaluated directly.

B Ensemble Size Ablation

We vary the ensemble size B from 1 to 100 and measure the standard deviation of the predicted quantile $\hat{q}_\tau(x)$ across five independent runs on the CIFAR-10 watermark set at $\tau = 0.999$ (FPR $= 0.1\%$). Table 13 reports the results. Prediction variance drops sharply from $B = 1$ to $B = 20$, continues to decrease up to $B = 50$, and plateaus beyond that point. We adopt $B = 50$ as the default, which offers a favorable trade-off between stability and computational cost. Since ensemble members are trained independently, scaling from $B = 50$ to any larger value is embarrassingly parallel and requires no inter-model communication.

Table 13: Effect of ensemble size B on quantile prediction stability (CIFAR-10, $\tau = 0.999$). Std. dev. computed over 5 independent runs.

Ensemble B	Std($\hat{q}_\tau$)	TPR @ FPR 0.1%
1	0.142	0.83 ± 0.05
10	0.068	0.87 ± 0.03
20	0.041	0.88 ± 0.03
50 (ours)	0.019	0.89 ± 0.03
100	0.016	0.89 ± 0.02

C Computational Resources

Table 14 summarizes the computational requirements of each component. The dominant cost is ensemble training (approximately 100 A100-hours for $B = 50$), which parallelizes trivially. At inference, t-error computation for 5,000 watermark samples takes roughly 10 minutes on a single A100, and ensemble quantile prediction adds less than 1 minute. In total, a single ownership query

(from raw images to verdict) completes in under 15 minutes, making the protocol practical for on-demand verification.

By comparison, explicit watermarking methods such as WDM [Peng et al., 2023] require re-training the diffusion model with the watermark objective, which typically takes as long as the original model training (on the order of days for the architectures we consider). MiO's Gaussian QR training is a one-time offline cost, and the owner model itself is never modified.

Table 14: Computational cost breakdown on CIFAR-10 (single NVIDIA A100).

Component	Time	Parameters
DDIM training (owner model)	$\sim$48 h	35.7M
Gaussian QR ($\times$1 model)	$\sim$2 h	11.2M
Gaussian QR ensemble (B=50)	$\sim$100 h†	560M total
Inference (per verification query)		
t-error scoring (5k samples)	$\sim$10 min	—
Ensemble prediction	$<$1 min	—
Two-point criteria	$<$1 s	—
Total query time	**$\sim$11 min**	—

† Embarrassingly parallel; wall-clock $\sim$2 h with 50 GPUs.

D Gaussian QR Training Details

Algorithm 3 Gaussian Quantile Regression Training

Require: Auxiliary non-member data $\mathcal{D}_{\text{aux}}$, ensemble size B, bootstrap ratio r
Ensure: Trained ensemble $\{f_\theta^{(b)}\}_{b=1}^B$
  for $b = 1$ to B **do**
    $\mathcal{D}_b \leftarrow$ BootstrapSample$(\mathcal{D}_{\text{aux}}, r)$ {80% with replacement}
    $\mathcal{D}_{\text{train}}, \mathcal{D}_{\text{val}} \leftarrow$ Split$(\mathcal{D}_b, 0.9)$
    Initialize $f_\theta^{(b)}$ with ResNet-18 backbone
    repeat
      for batch $(x, s(x), \phi(x))$ in $\mathcal{D}_{\text{train}}$ **do**
        $y \leftarrow \log(1 + s(x))$
        $(\mu, \log\sigma) \leftarrow f_\theta^{(b)}(x, \phi(x))$
        $\mathcal{L} \leftarrow (y - \mu)^2/(2\sigma^2) + \log\sigma$
        Update θ via gradient descent
      end for
    until early stopping (patience = 10)
  end for
  return $\{f_\theta^{(b)}\}_{b=1}^B$

E Threat Model Boundary: Distillation

An adversary who generates a large synthetic dataset from the stolen model and trains a new model from scratch effectively performs *knowledge distillation*. Because the student is initialised independently and never sees the owner's training data directly, per-example memorisation of $\mathcal{W}$ is not expected to transfer.

Experiment. We generate 50,000 synthetic CIFAR-10 images from Model A via 10-step DDIM sampling and train a student model from scratch (100,000 iterations, batch size 128, fresh random initialisation). Table 15 reports t-error statistics on $\mathcal{W}$.

The student's mean t-error on $\mathcal{W}$ is 106.9, compared to 28.7 for the owner—a 3.72$\times$ gap. The verification protocol correctly returns FAIL: the consistency criterion (C1) rejects because the

Table 15: Distillation boundary experiment on CIFAR-10. The student model, trained on synthetic data from the owner, exhibits 3.72× higher t-error than the owner and fails the consistency criterion.

Model	Mean t-error	Std	Ratio vs. Owner
Model A (owner)	28.71	10.85	1.0×
Student (distilled)	106.90	36.14	3.72×
ddpm-cifar10 (reference)	703.84	38.50	24.5×
Random UNet (lower bound)	3037.31	65.75	105.8×

student's score distribution differs significantly from the owner's ($p \approx 0$), indicating the student does not reproduce the owner's per-example memorisation pattern.

This outcome is expected and represents a principled boundary of the method rather than a failure mode. Distillation with fresh initialisation approximates full retraining: generating a sufficiently large synthetic dataset from a diffusion model and training a new model from scratch approaches the computational cost of the original training procedure, undermining the economic motivation for model theft. Critically, the protocol does not produce a false positive—it correctly reports that ownership cannot be established, which is the far more important guarantee.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and introduction (Section 1) claim Cohen's $d > 18$ and error ratios $> 19\times$ across four datasets, with robustness under MMD fine-tuning. These are directly supported by Table 2 and Section 5.3. Limitations regarding white-box access and full retraining are stated in Section 6.

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Section 6 explicitly discusses five scope conditions: white-box access requirement, watermark retention, baseline availability, the Gaussian assumption, and what the work does not address (API-only verification, full retraining adversaries, memorization erasure).

   Guidelines:

   - The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The Gaussian QR formulation (Section 4.3) states its distributional assumption explicitly ($y \sim \mathcal{N}(\mu(x), \sigma^2(x))$) and derives the closed-form quantile threshold. The combinatorial bound on watermark guessing (Section 3) is a counting argument with all assumptions stated.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Table 1 and Appendix A provide all hyperparameters for DDIM training, Gaussian QR, and MMD fine-tuning. Dataset splits are deterministic with fixed seeds. Public baseline checkpoints are identified by their HuggingFace model IDs. Algorithm 1 and Algorithm 2 formalize the procedures.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

(d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymized code repository containing the MiO verification protocol, training scripts, and evaluation pipeline will be provided as supplementary material with this submission. See Section 5.1 for setup instructions and Appendix A for hyperparameters.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 5.1 specifies datasets, split sizes, seeds, and baseline models. Table 1 summarizes key hyperparameters. Appendix A provides full architectural details (U-Net channels, attention resolutions, dropout), optimizer settings (AdamW, learning rate, weight decay), and MMD fine-tuning configuration.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Table 3 reports 95% bootstrap confidence intervals (200 iterations) for TPR at fixed FPR. The verification protocol itself is grounded in hypothesis testing (p-values, Cohen's d), and all p-values for baseline separation are $< 10^{-100}$ (Table 2).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix C (Table 14) provides a detailed computational cost breakdown on NVIDIA A100 GPUs, including DDIM training time ($\sim$48h), Gaussian QR ensemble cost ($\sim$100 A100-hours, parallelizable to $\sim$2h wall-clock), and per-query inference time ($\sim$11 min).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

Answer: [Yes]

Justification: The research uses publicly available datasets (CIFAR-10/100, STL-10, CelebA) and public pretrained models from HuggingFace. The work addresses IP protection, a legitimate research goal. No human subjects are involved.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.

- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: This work enables IP holders to verify model ownership with objective statistical evidence (positive impact). Potential risks include false ownership claims and retrospective selection of evidence sets; these are mitigated by requiring pre-registered watermarks with hash commitments, comparison against multiple public baselines, and principled significance thresholds (Section 4.4). Section 6 further discusses scope conditions.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: This paper proposes a verification protocol, not a new generative model or dataset. No new models or data with misuse risk are released.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All datasets (CIFAR-10/100, STL-10, CelebA) and pretrained models (HuggingFace DDPM checkpoints) are cited with their original publications in the references. These are standard research benchmarks with permissive research-use terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not release new datasets, pretrained models, or code packages as assets.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This work does not involve crowdsourcing or human subjects research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.

- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This work does not involve human subjects research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not used as a component of the core method. The method relies on diffusion model reconstruction error and Gaussian quantile regression, with no LLM involvement.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.